

FINAL PRESENTATION · MGTA 495

Gymshark AI Influencer Marketing Pipeline

End-to-end creator discovery, scoring, outreach & campaign
analytics



Today's discussion

01 Current business process — how Gymshark does it today

02 AI-supported workflow — full swimlane process map

03 Data & evidence used by the AI

04 Human review, edit, rejection & approval steps


05 Time, cost & quality estimates

06 Face-validity check

07 AI performance — actual test results


08 What we'd improve with more time

Today the creator pipeline is manual, sequential, and slow



Current manual workflow · ≈ 80–135 min per creator · 4 handoffs · 2–3 day turnaround.

The AI-supported workflow spans five phases and four roles



● AI coworker task ● Human + AI collaboration ● Manager review / approval ● Human task ● Final output

Full BPMN swimlane — five phases, four roles. AI performs the work; humans review and hold two approval gates.

The AI runs on six structured datasets and a grounded knowledge base

Structured data file	Contents
creators.csv	12 YouTube prospects · GFS · brand safety · deal type
comment_sentiment.csv	VADER NLP · compound score · purchase intent · negative flags
crm_records.csv	Pipeline stage · account manager · channel · priority
hypeauditor_intake_log.csv	Fraud risk · fake-follower % · audience quality
existing_partners.csv	Current partners — kept separate from prospects
campaign_performance.csv	ROI · revenue · reach · UTM clicks · promo redemptions

RAG knowledge base — documents indexed

- Gymshark brand guidelines
- Outreach playbook & banned-phrases list
- Deal-type policy & product catalog (ADAPT / TRAINING / LIFTING / RUNNING)

Retrieval method

- TF-IDF index · top-3 chunks per query
- Title, section, relevance score + 300-char snippet shown in the UI

Why it matters

- Every draft & brief traces to a specific brand document
- Fix a wrong output by editing the knowledge base — not by re-prompting

Humans approve every external action — the AI never decides alone

Pipeline step	Approval	What the AI does	What the human does
Creator scoring	Yes — Analyst	Calculates GFS · auto-disqualifies red / clickbait / competitor	Reviews ranked leaderboard · overrides amber flags · moves creators forward
Outreach draft	Yes — Analyst	Retrieves 3 RAG chunks · generates personalised email	Checks banned phrases · adjusts tone and deal type · submits
Outreach approval	Yes — Manager	Nothing at this step	Manager approves, requests edit with reason, or rejects — email cannot leave without this
Product brief	Yes — Manager	Selects product line · generates promo code · writes brief	Manager can override product line · approves or regenerates before sending
Campaign analytics	Informational	Calculates ROI · flags top / bottom performers · outputs dashboard	Reviews data · makes the final renewal or archive decision

No outreach leaves without analyst submission + manager approval; no brief is sent without manager sign-off. The AI accelerates — it does not decide.

Time, Cost & Quality Impact

The pipeline saves time by automating repetitive evidence-gathering while keeping human review in the loop for approval and risk decisions.

EFFICIENCY CASE

60–65%

TIME REDUCTION

Plausible savings when initial screening and enrichment are automated.

£37

SAVED / CREATOR

Manual: £53.50
AI-assisted: £16.50

~£1,850

MONTHLY IMPACT

saved at 50 creators/month

Even under a **conservative 20–30%** total reduction, the gain remains meaningful because the highest-friction research loop is automated.

QUALITY GUARDRAILS

Savings do not come from removing judgment.

AI HANDLES

search, enrichment, scoring, and draft generation

HUMANS RETAIN

final approval for outreach and risk cases

UNKNOWN RISK

escalated when comments are insufficient

DEEPEVAL CHECKS

faithfulness, relevance, business quality, and face validity

HUMAN-IN-THE-LOOP

WORKFLOW SHIFT

BEFORE

Manual vetting required repeated searching, watching, copying notes, scoring, and drafting per creator.

45–90 MIN / CREATOR

AFTER

One script run produces scored, CRM-ready records; reviewers focus on edge cases and decisions.

SECONDS FOR AUTOMATED PASS

KEY POINT

The value is faster, more consistent first-pass screening — **not replacing the manager's final judgment.**

The output holds up to formula design, grounding, and external benchmarks

Scoring formula design

- Authenticity (30%) and comment sentiment (25%) dominate
- Red brand safety –2.0 · clickbait –1.5 · competitor –1.0
- Penalties applied before any subjective judgment

RAG grounding

- Every draft cites the brand documents it drew from
- Wrong output → fix the knowledge-base chunk, not the prompt
- Traceable, auditable, correctable

External benchmark

- HypeAuditor 2025: manual vetting = 45–90 min / creator
- Automated enrichment = under 10 min
- Our pipeline sits within this range

Direction is correct

- AI removes the repetitive steps — search, read, score, log
- Human review stays roughly constant; the bottleneck shifts to structured decisions
- Consistent with how AI augments knowledge workflows

Honest caveat

- Campaign performance and HypeAuditor data are simulated
- Conservative estimate: a 20–30% time reduction is still meaningful
- We report both the optimistic and conservative range

AI Evaluation Framework & Key Findings

DeepEval turns LLM output quality into repeatable tests for outreach, brand-fit, and product briefs.

TEST SCOPE

13

AUTOMATED TEST CASES

3 PIPELINE OUTPUTS

7

Outreach drafts

3

Brand-fit analyses

3

Product briefs

9/13

PASSED

100% faithfulness + relevance

No hallucination pattern detected

DEEPEVAL CHECKS

Judge model

API-GPT-OSS-120B

Workflow

API-LLAMA-4-SCOUT

Faithfulness

Grounded in creator data?

Relevance

Answered the requested task?

Business Quality

Followed Gymshark rules?

Face Validity

Ready for manager approval?

DeepEval gave the team a repeatable scoring harness: every prompt, model, or rule change can be re-tested against the same quality bar.

FINDINGS & FIXES

MAIN FINDING

Failures were quality / judgment issues, not factual errors.

ISSUE → FIX

Issue	Fix Applied
Tone sycophancy	Require factual first-line observation
High-risk hedging	Force human-review sentence
Generic briefs	Inject niche into KEY MESSAGE
Face Validity	0.40 → 1.00

TAKEAWAY

Automated AI testing converts subjective review into measurable, repeatable prompt improvements.

Domain knowledge is the biggest gap to close next

More domain knowledge — the biggest gap

- Many assumptions made without real insider context
- How the team evaluates creators, structures deals, and defines renewal
- A real partnerships manager would reshape the scoring formula, RAG docs & brief templates
- Architecture is sound — calibration needs real-world grounding

Real creator data

- Campaign performance and HypeAuditor data are simulated today
- Real YouTube Analytics — engagement, retention, demographics
- Live HypeAuditor responses for actual fraud-risk scores
- Simulated data tests structure, not whether predictions hold

Closed-loop feedback

- **Implemented** — outreach outcomes feed back to scoring via an LLM-based loop
- Next step: replace the LLM judge with a **trained statistical model** for lower latency, lower cost, and reproducible scores
- A labelled outcomes dataset from live campaigns would enable this transition

Finer-grained testing

- LLM-dependent tests need mocking or a dedicated API environment
- End-to-end integration test: discovery → scoring → outreach → brief
- A fixed creator dataset to catch cross-module regressions early

IN SUMMARY

The system gathers the data;
the team keeps the decisions.

The AI scores, writes, and computes. Brand expertise and every external action stay with the team. Domain knowledge is the next frontier.

234

tests passing

~63%

time reduction

£1,850

saved / month

5

pipeline phases

2

manager gates